Literature Review

Data Visualisation for Economic Crisis Prediction

University of Essex

MSc Artificial Intelligence

Module 6: Research Methods and Professional Practice

Pavlos Papachristos (student id: 12697110)

Literature Review

Data Visualisation for Economic Crisis Prediction

1. Introduction

Financial crises tend to break out abruptly, even though in many cases the recurring warning signs made them to look obvious retrospectively. The Great Depression, the 2008 global crisis and the financial disruption following COVID-19 have showed how pressures can build gradually before suddenly becoming unmanageable. Despite improvements in economic monitoring, the sheer volume and velocity of macro-financial data make early detection difficult.

Visualisation tools have started to ease part of this burden. By converting large datasets into clearer visual forms, have allowed irregularities and emerging patterns to be spotted earlier than would be possible from raw tables or lengthy briefing documents. This review considers how these visual tools contribute to crisis prediction, what they reveal that numerical approaches may obscure, and where their limitations remain. It draws on both traditional statistical work and recent machine-learning research, along with the visual frameworks used in institutional settings.

2. Methodology Applied

The review draws on literature published between 1998 and 2025, identified through searches of Google Scholar, Web of Science, EconLit and SSRN, supplemented with IMF, BIS, ECB and Federal Reserve publications. Studies were included when they presented empirical evidence or contained a clear link between visualisation and crisis prediction. The literature spans early-warning work such as Kaminsky, Lizondo and Reinhart (1998), long-run historical work like Reinhart and Rogoff (2009), institutional frameworks from the IMF and ESRB, and recent machine-learning research (Beutel, List and von Schweinitz, 2019; Molnar, 2020).

3. Data Visualisation Evolution in Economics

Economists have relied on visual aids for a long time, but their role has expanded as the amount of financial data has increased. The early warning systems that Kaminsky, Lizondo and

Reinhart (1998) studied were quite simple by today's standards. They analysed 15 indicators across 20 countries spanning 1970-1995, using threshold-based signal extraction. Their approach correctly identified 70% of crises within 24 months but generated substantial false positives.

With increasing global financial integration, researchers shifted toward more complex techniques. Frankel and Saravelos (2012) used 31 indicators across 60 countries, combining visual examination with econometric methods during the 2008-09 crisis. This richer framework demands greater technical expertise, making interpretation more difficult for policymakers.

With increasing global financial integration in the 2000s, researchers shifted toward more complex techniques. Drawing on 31 indicators across 60 countries, Frankel and Saravelos (2012) found that combining visual examination with econometric methods helped them reach more precise conclusions during the 2008-09 crisis. This framework demands a greater technical expertise that makes it more difficult for analysis and interpretation by the policymakers.

There is an apparent trade-off: simple visual methods are intuitive but may overlook deeper interactions, while complex multivariate systems capture nuance at the cost of transparency.

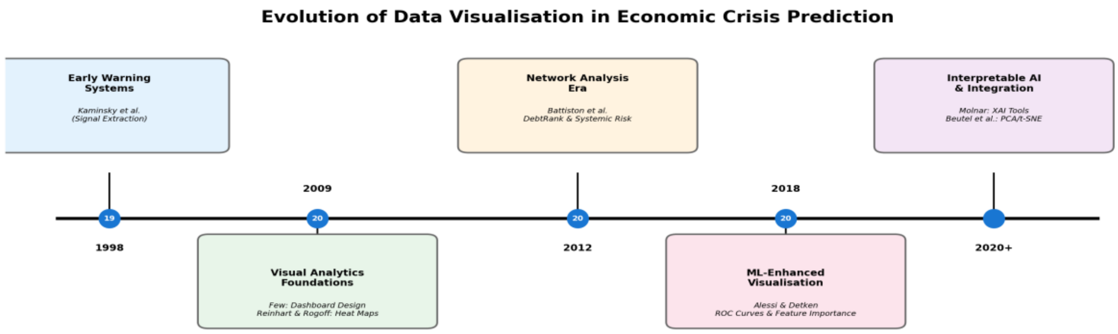

Figure 1: Evolution of visualisation approaches in economic crisis prediction (adapted from Kaminsky, Lizondo and Reinhart, 1998; Few, 2009; Battiston et al., 2012; Alessi and Detken, 2018; Molnar, 2020)

4. The Main Visualisation Techniques

4.1 Traditional Approaches

Time-series charts remain the basic tool for tracking financial indicators. Claessens and Kose (2014) relied heavily on them when studying indicators such as the credit-to-GDP ratio. Their work showed that simple patterns—slopes steepening, trends breaking—often give early clues that conditions are shifting.

Heat maps serve a different purpose. Reinhart and Rogoff (2009) used them to compress centuries of cross-country crisis data into a single format. Colour gradients highlight shared vulnerabilities that are difficult to spot in individual time-series plots.

4.2 Interactive Tools

Few (2009) argued strongly for interactive dashboards, emphasising that the ability to adjust filters, time windows or country selections allows analysts to explore data more freely. This matters because financial vulnerabilities often emerge in specific regions or sectors rather than in aggregate indicators.

Network visualisations add another dimension. Battiston et al. (2012) showed that mapping institutional linkages reveals risk propagation channels that are not visible in standard macro-financial indicators. Identifying which institutions are in the centre of dense webs of exposures offers a clearer sense of how shocks might spread across the system.

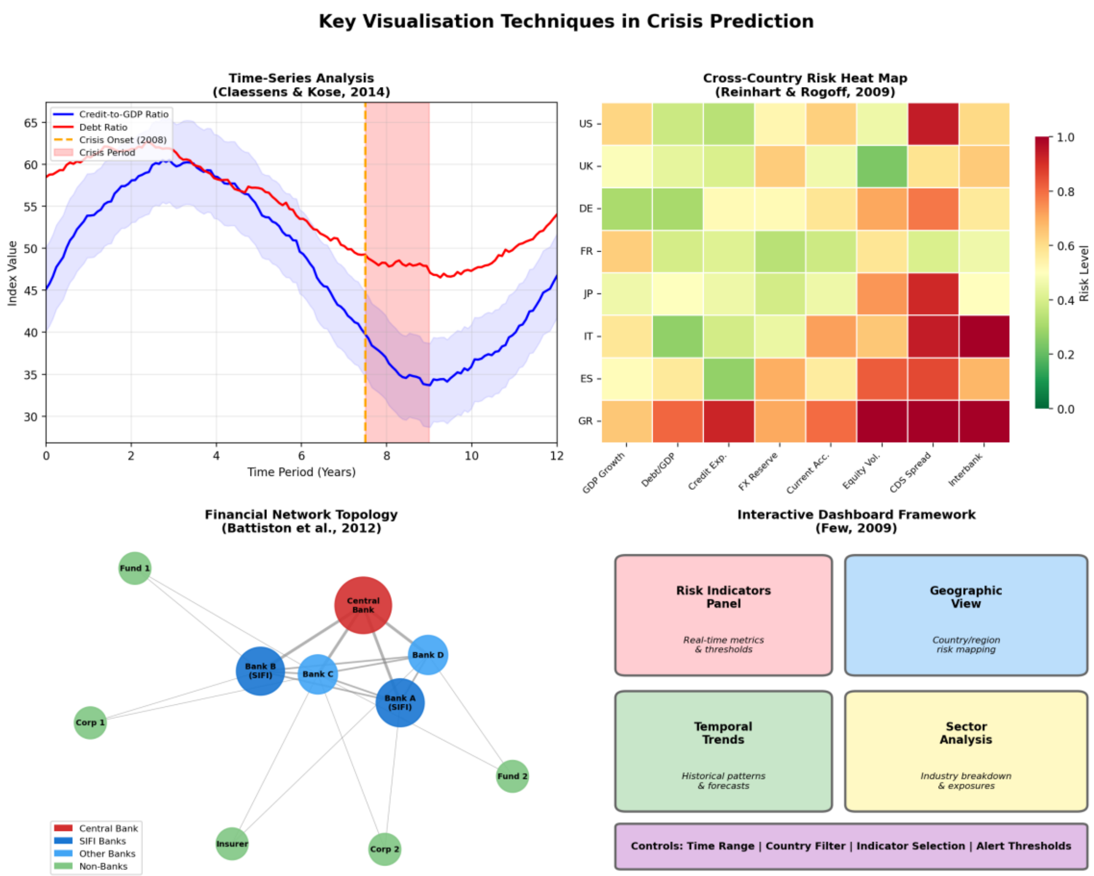

Figure 2: Main visualisation techniques used in crisis prediction research (based on Claessens and Kose, 2014; Reinhart and Rogoff, 2009; Battiston et al., 2012; Few, 2009)

4.3 Critical Evaluation

The progression from static techniques to interactive formats represents genuine development, though it's unclear whether the added complexity consistently improves outcomes. Few (2009) makes a strong case for dashboards but provides little empirical support showing that interactivity enhances decision-making relative to high-quality static displays. Network visuals add insight but demand greater technical confidence from users.

5. Machine Learning and Visual Analysis

Machine-learning models offer new possibilities for predictive work, but they depend heavily on visual tools to make their outputs interpretable. Alessi and Detken (2018) showed that ROC curves and feature-importance visuals can make even standard models far easier to evaluate.

Beutel, List and von Schweinitz (2019) applied ML techniques to a large global dataset covering 60 countries from 1980 to 2017. They used dimensionality-reduction tools like PCA and t-SNE to display high-dimensional patterns in two or three dimensions. Their conclusions were cautious: machine-learning models only modestly improved on a simple logit benchmark (AUC 0.81 vs. 0.78). Their visuals also revealed that the models were mostly detecting well-known signals—credit booms, external imbalances—rather than uncovering new predictors.

Molnar (2020) emphasises the potential of interpretability tools such as partial-dependence plots. These can clarify how models behave, but their value depends on whether the underlying model has found genuinely new relationships.

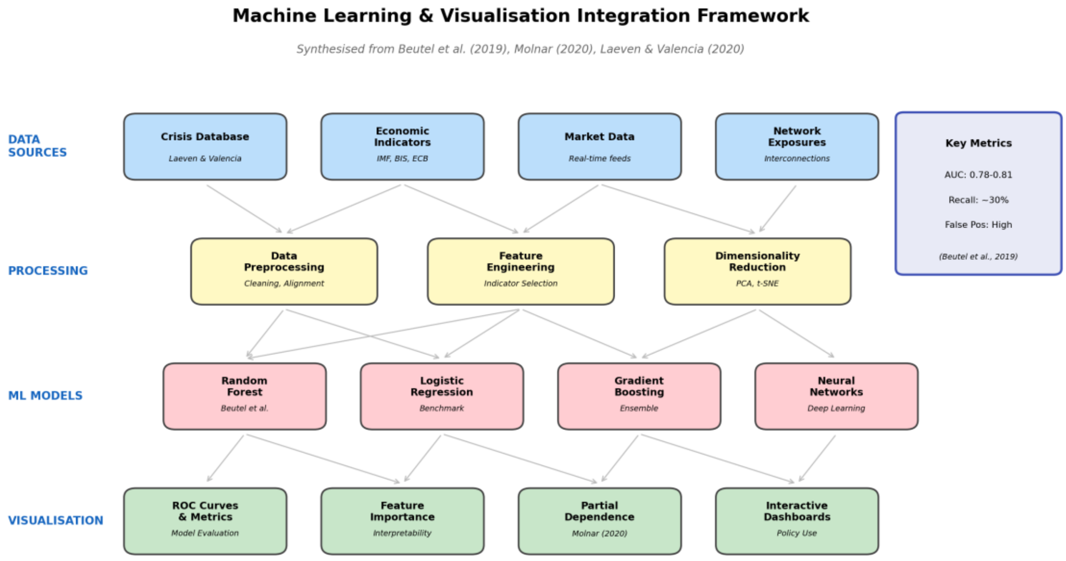

Figure 3: Machine learning and visualisation integration in crisis prediction frameworks (compiled from Beutel, List and von Schweinitz, 2019; Molnar, 2020; Laeven and Valencia, 2020)

5.1 Critical Evaluation

Comparing Alessi and Detken's euro area work with Beutel et al.'s global analysis reveals something important. Studies focusing on a single region tend to do better, probably because economic structures and institutions are more consistent. But what works in Europe may not transfer to Asia or Latin America.

When Beutel, List and von Schweinitz (2019) tested models across 60 countries, the improvements over basic methods were modest. This raises concerns whether interpretability techniques are showing us something new or just giving us a fancier way to see familiar warning signs. The case for complex machine-learning methods seems strongest when dealing with genuinely messy, high-dimensional data where relationships don't follow simple patterns.

6. Real Applications in the Field

Regulatory and policy maker institutions have developed visual systems that differ substantially in scope and design. The IMF's Financial Soundness Indicators dashboard (Čihák, 2007) aggregates more than 40 indicators across nearly 200 countries, using time-series charts, heat maps and cross-country panels to provide a broad and consistent overview.

The ESRB's dashboard (Constâncio, 2017) takes a more focused approach, concentrating on the euro-area and integrating network diagrams that show how risks might spread through interconnected institutions. This depth-oriented design matches the ESRB's role monitoring a tightly linked regional financial system.

Although both dashboards are widely used, neither provides convincing evidence of predictive success. Constâncio's (2017) suggestion that ESRB tools "helped spot problems" is not supported by formal evaluation. IMF assessments rely heavily on usage metrics, which indicate demand but not improved forecasting. Without structured validation, it is difficult to determine whether these dashboards genuinely enhance crisis detection or simply present familiar data more neatly.

6.1 Critical Evaluation

The tools Constâncio (2017) and the ESRB describe are built mainly for usability—they help people make decisions quickly. Statistical precision comes second. That's a reasonable trade-off when spotting emerging problems in real time, but it means these systems get validated differently than academic models. The focus is on whether they make sense to users and whether stakeholders find them credible, rather than rigorous testing against holdout data.
The IMF's vulnerability assessments rely heavily on expert judgment informed by visual displays rather than formal predictive algorithms. This raises questions: can someone else reproduce the same assessment? Are these tools making forecasts more accurate, or mainly just a better way of organising analysis that would have happened anyway?

7. Problems and Limitations

Rose and Spiegel (2011) highlight the fundamental difficulty of predicting rare events: limited samples, shifting regimes and structural breaks restrict what any model or visualisation can achieve. Visuals may clarify patterns, but they cannot compensate for weak data.

As Taleb (2007) notes, people tend to read structure into information even when no real pattern exists. Few (2009) adds that interactive dashboards can intensify this habit—continual filtering can make chance fluctuations look meaningful. In effect, this becomes a visual form of p-hacking (manipulating data analysis until finding statistically significant results even if spurious).

Network diagrams introduce other risks. Battiston et al. (2012) show that visually dominant clusters can create a sense of systemic importance not fully supported by centrality metrics.

There's little empirical work examining whether visual tools raise false-positive rates or distort judgement.

7.1 Conflicting Evidence and Unresolved Debates

Several debates remain unsettled. One question concerns whether complex models outperform simpler indicators. Frankel and Saravelos (2012) show that standard measures performed convincingly during the 2008-09 crisis, whereas Alessi and Detken (2018) and Beutel, List and von Schweinitz (2019) report only modest improvements from machine-learning methods.

Another debate centres on interpretability. Molnar (2020) and Few (2009) advocate that visual aids make sophisticated algorithms more accessible, while Taleb (2007) warns that reducing nonlinear dynamics to tidy graphics may distort reality.

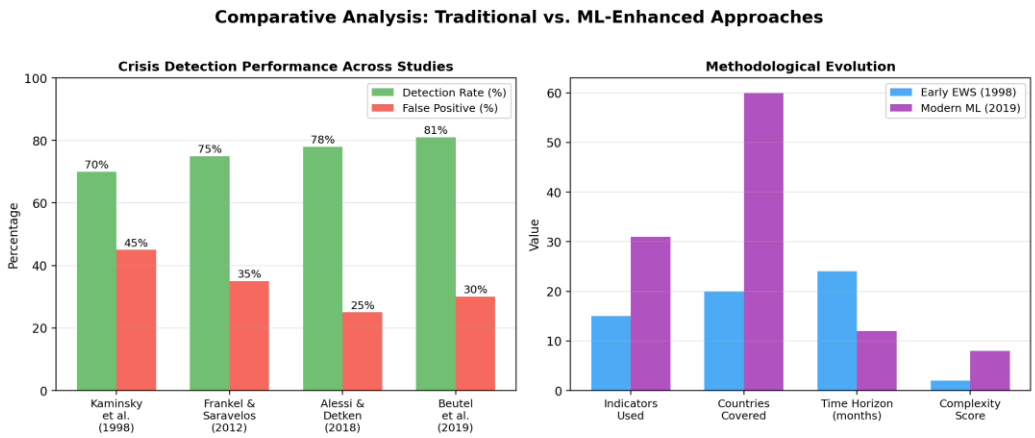

Figure 4: Comparative analysis of crisis detection performance across key studies (compiled from Frankel and Saravelos, 2012; Alessi and Detken, 2018; Beutel, List and von Schweinitz, 2019)

8. Research Approaches and Methods

Research mixes quantitative modelling with qualitative judgement. Laeven and Valencia (2020) show how much preliminary work is required: crises must be identified through several criteria, definitions kept consistent across countries, and indicators aligned to avoid timing mismatches.

Validation practices differ markedly. Alessi and Detken (2018) adopt a temporal split, training models on pre-2005 data and testing on the following decade. Beutel, List and von Schweinitz (2019) add an important caution: models with strong AUC scores may still perform poorly at practical thresholds, identifying only around 30% of true pre-crisis periods.

9. Gaps and Future Directions

Several possibilities remain unexplored. Emerging data sources such as social-media activity, offer a potential insight, yet there is no established method for integrating them visually with traditional macro-financial indicators. Their structure differs enough that standard formats often prove inadequate. The temporal dynamics of crises also require further work. Most visuals still rely on either static views or simple timelines, while real crises evolve through feedback loops and shifts in behaviour that are hard to depict clearly.

More broadly, despite extensive design guidance, there is little evidence on which specific visual formats genuinely aid crisis detection or best support policymakers.

10. Concluding Thoughts

Visualisation has become increasingly important because it helps analysts navigate large datasets and notice patterns that would otherwise remain hidden. However, the challenges that undermine crisis prediction still apply. The challenges however, that undermine crisis prediction more generally still apply. Crises remain rare, data is often noisy or incomplete and patterns can easily be misleading. Machine-learning techniques, paired with visual tools, offer new possibilities but cannot overcome these structural issues on their own. Their value increases only when supported by sound theory, careful data handling and realistic validation.

As financial systems become more interconnected, the need for clear, reliable visual tools will grow even further. The main question is not whether visualisation will play a role but whether the tools can evolve quickly enough to make sense of systems that are themselves becoming more complex.

References

Alessi, L. and Detken, C. (2018) 'Identifying excessive credit growth and leverage', Journal of Financial Stability, 35, pp. 215-225.

Battiston, S., Puliga, M., Kaushik, R., Tasca, P. and Caldarelli, G. (2012) 'DebtRank: Too central to fail? Financial networks, the FED and systemic risk', Scientific Reports, 2(1), pp. 1-6.

Beutel, J., List, S. and von Schweinitz, G. (2019) 'Does machine learning help to predict financial crises?', Journal of Financial Stability, 45, 100693.

Claessens, S. and Kose, M.A. (2014) 'Financial crises explanations, types, and implications', in Claessens, S. and Kose, M.A. (eds.) Financial Crises: Causes, Consequences, and Policy Responses. Washington: International Monetary Fund, pp. 3-58.

Constâncio, V. (2017) 'Macroprudential policy in a changing financial system', Speech at the Second ECB Macroprudential Policy and Research Conference, Frankfurt, 11 May. Available at: https://www.ecb.europa.eu (Accessed: 20 November 2025).

Few, S. (2009) Now You See It: Simple Visualization Techniques for Quantitative Analysis. Oakland: Analytics Press.

Frankel, J. and Saravelos, G. (2012) 'Can leading indicators assess country vulnerability? Evidence from the 2008-09 global financial crisis', Journal of International Economics, 87(2), pp. 216-231.

Kaminsky, G., Lizondo, S. and Reinhart, C.M. (1998) 'Leading indicators of currency crises', IMF Staff Papers, 45(1), pp. 1-48.

Laeven, L. and Valencia, F. (2020) 'Systemic banking crises database II', IMF Economic Review, 68(2), pp. 307-361.

Molnar, C. (2020) Interpretable Machine Learning: A Guide for Making Black Box Models Explainable. Available at: https://christophm.github.io/interpretable-ml-book/ (Accessed: 20 November 2025).

Reinhart, C.M. and Rogoff, K.S. (2009) This Time Is Different: Eight Centuries of Financial Folly. Princeton: Princeton University Press.

Rose, A.K. and Spiegel, M.M. (2011) 'Cross-country causes and consequences of the crisis: An update', European Economic Review, 55(3), pp. 309-324.

Taleb, N.N. (2007) The Black Swan: The Impact of the Highly Improbable. New York: Random House.